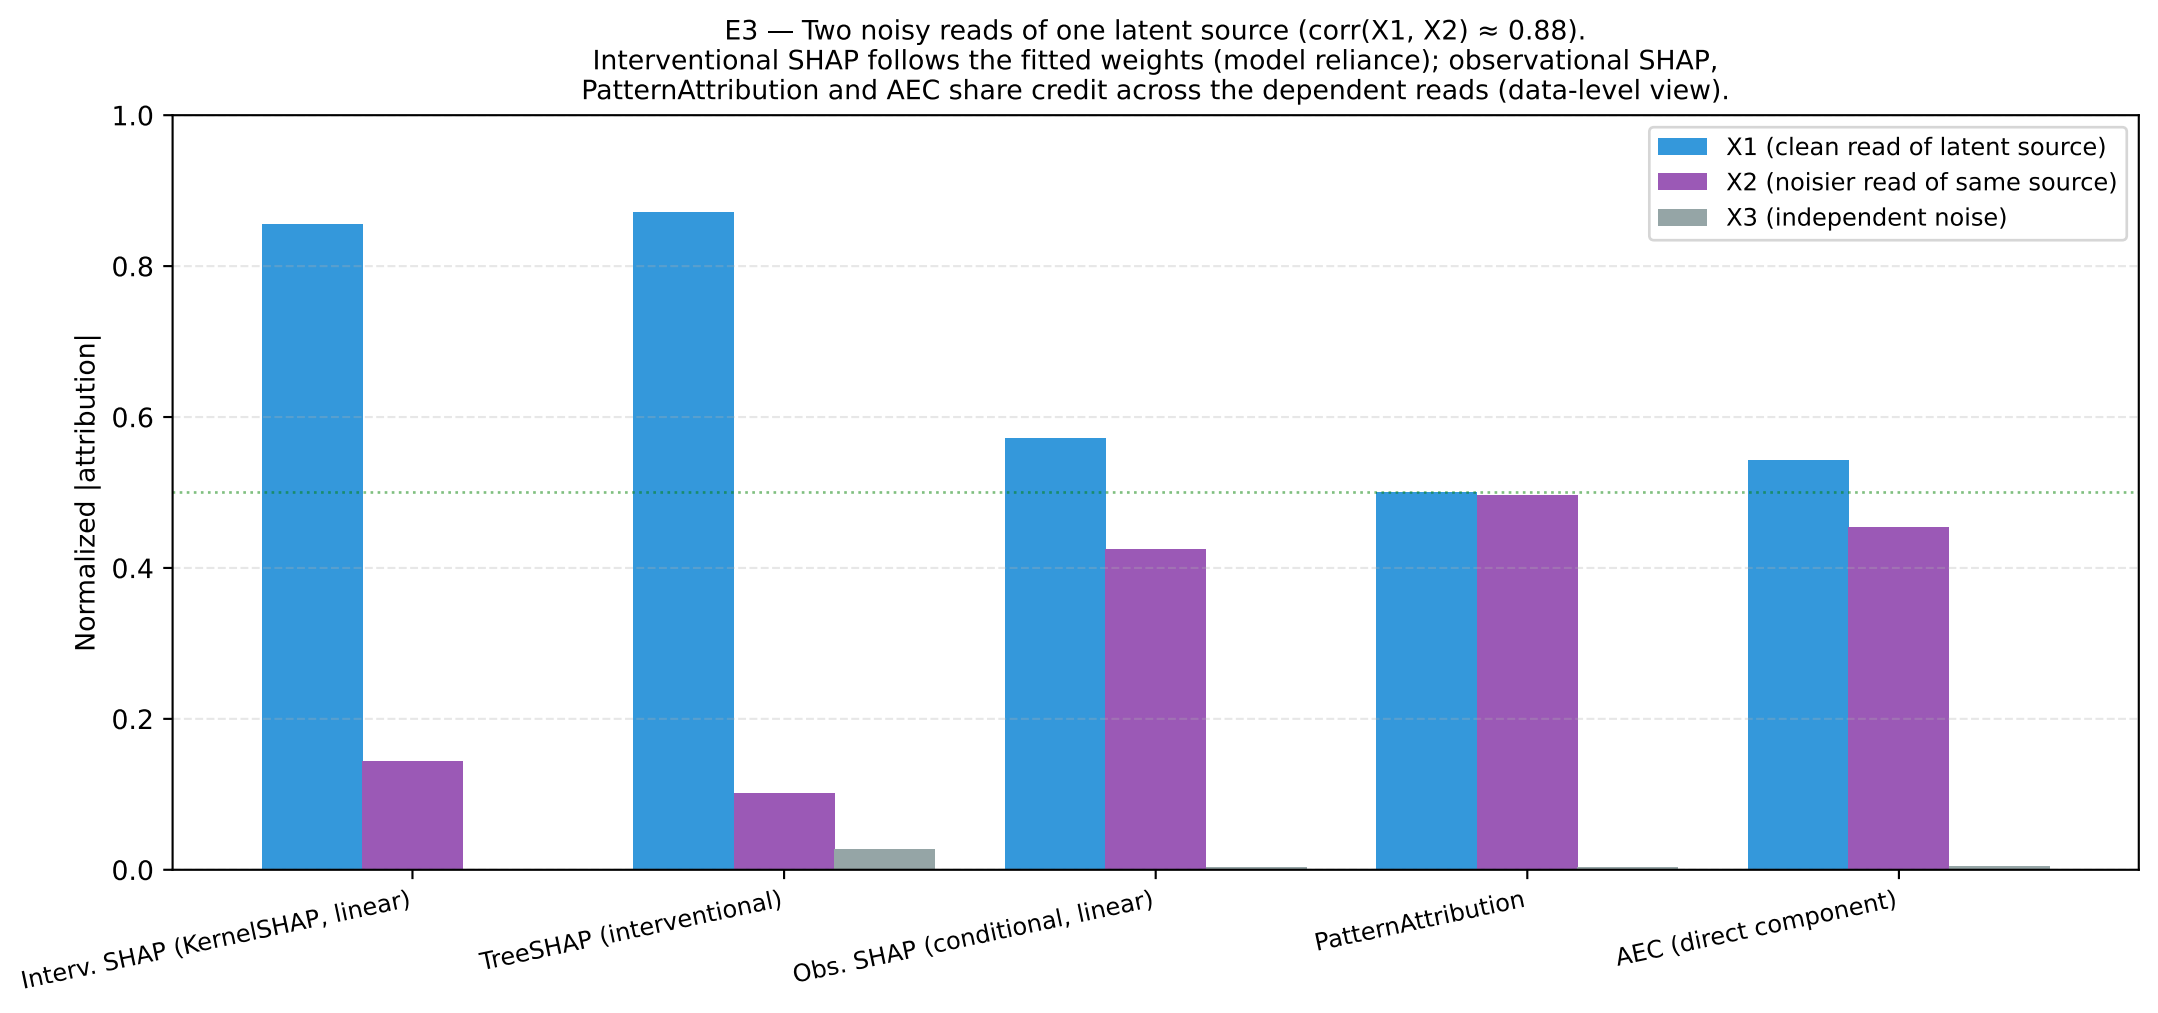
E3 — Two noisy reads of one latent source (corr(X1, X2) ≈ 0.88).
Interventional SHAP follows the fitted weights (model reliance); observational SHAP,
PatternAttribution and AEC share credit across the dependent reads (data-level view).
Normalized |attribution|
1.0
0.8
0.6
0.4
0.2
0.0
X1 (clean read of latent source)
X2 (noisier read of same source)
X3 (independent noise)
Interv. SHAP (KernelSHAP, linear)
TreeSHAP (interventional)
Obs. SHAP (conditional, linear)
PatternAttribution
AEC (direct component)